

DSR - Introduction to Reinforcement Learning

DSR batch # 17 - February 2019

About me

- DSR participant - batch #15
- Msc in Engineering and Applied Mathematics
- Now working at Ecole Polytechnique in Paris as a Research Engineer in Machine Learning and PhD student (in Reinforcement Learning)
- LinkedIn: <https://www.linkedin.com/in/alicemartindonati/>
- Email: alice.martindonati.pro@gmail.com

Course outline

- Introduction: Reinforcement Learning in the Machine Learning Field
- Part I: Key concepts in Reinforcement Learning
- Part II: Solving a MDP with Dynamic Programming
- **Implementation of DP algorithms in Python**
- Part IV: Tabular Q-learning
- **Implementation of tabular-Q learning in Python**
- Part V: Introduction to Deep Q networks
- Part VI: Introduction to policy gradient algorithms
- Conclusion: Future of RL and challenges


The Reinforcement Learning problem

At each timestep t , the Agent:

- receives an observation O_t and reward R_t from the environment
- performs an action A_t

The environment:

- receives action A_t ,
- emits observations O_{t+1} and scalar reward R_{t+1}



RL: A decision making problem where an agent interacts sequentially in an environment

Goal: find the sequential set of actions that maximizes his total reward

Machine Learning and Reinforcement Learning

Supervised Learning

Labeled Dataset (X,y)
Learn Relation
between X and y
 $y=f(X)$
Example: Image
Classification

Unsupervised Learning

Unlabelled Dataset X
Learn Hidden Pattern
in the Dataset
Example: Clustering

Reinforcement Learning

**Agent Learns by
interacting with an
environment and
receiving a signal
from this
environment**

**-Trial and Error
Process**

Characteristics of Reinforcement Learning

- There is no supervisor, only a reward final
- Feedback is delayed, not instantaneous
- Time really matters: sequential decision making problem
- Active learning process: Agent's actions influence the subsequent data it receives (the reward he gets, his state).

Example : The frozen Lake

S	F	F	F
F	H	F	H
F	F	F	H
H	F	F	G

An agent has to go from S to G in a frozen lake without falling into the holes.

Actions: move right, left, up, down.

The agent receives +1 if he reaches the goal G, 0 if he falls into a hole.

The episode ends when the agent falls into a hole or reaches the goal

Examples of sequential decision making problems

Fly stunt manoeuvres with an helicopter

Make a humanoid robot walk

Beat humans at Games like Go or BackGammon

Play Atari Games

Energy: Control a power station

Finance: Optimise an investment portfolio

Medicine: Optimize drug dosages

Key Elements in a RL problem

OBSERVATIONS:

- History H_t : sequence of past observations
- State S_t : function of H_t


SET OF ACTIONS A

REWARD = function of state S_t and action A_t

DISCOUNT RATE γ : penalisation of the future

MODEL OF THE ENVIRONNEMENT

TERMINAL STATE S_T OR TIME HORIZON H



Key Elements of a RL problem - History, State and the Markov Property

- History: sequence of past observations

$$H_t = (O_1, R_1, A_1), \dots, (O_t, R_t, A_t)$$

- The state is a function of the history :

$$S_t = f(H_t)$$

- Markov Property : “The future is independent of the past given the present”

$$P(S_{t+1} | S_t) = P(S_{t+1} | S_t, \dots, S_1)$$

- Once the state is known, the history can be thrown away

Key Elements of a RL problem - The reward

- A reward R_t is a scalar feedback signal
 - It indicates how well the agent is going at step t
 - The agent job is to maximise **cumulative rewards**
-
- Reinforcement Learning is based on the reward hypothesis: all goals can be described by the maximisation of expected cumulative reward

Key Elements of a RL problem - The Return

- The return G_t is the total discounted reward from time step t

$$G_t = r_{t+1} + \gamma r_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

- The discount rate γ is the present value of future rewards: $\gamma \in [0, 1]$

Key Elements of a RL problem - The Discounted rate

Why using a discounted rate? **The discount rate penalizes the future**

- Future is more uncertain
- Human Preference for immediate rewards
- If reward = money, discounted rate corresponds to actualisation rate
- Mathematical convergence for environnement with no terminal states:

$$G_t = r_{t+1} + \gamma r_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

Key Elements of a RL problem - The Agent's Policy


- Model the agent's behaviour
- Mapping from States to Actions

- Deterministic Policy:

$$a = \pi(s)$$

- Stochastic Policy:

$$\pi(a|s) = P(A_t = a|S_t = s)$$



Key Elements of a RL problem - The Value Functions

- The Value Function is a prediction of future reward
- Used to evaluate the Goodness / Badness of each state
- And therefore to select between actions

2 Value-Functions:

- **Value Function $V(s)$: Measures how good is being in a state s**
- **Action Value Function $Q(s,a)$: Measures how good is being in a state s and taking an action a**

Value-Functions computations are the core of most reinforcement learning algorithms

Key Elements of a RL problem - The model of the environment

A model is the agent's representation of the environment

A model is made of 2 components:

- **Transition Model P: predicts the next state**
- **Reward Model R: predicts the next immediate reward**

$$p(s'|s, a) = P \left[S_{t+1} = s' \mid S_t = s, A_t = a \right]$$

$$r(a, s) = \mathbb{E} [r_{t+1} \mid S_t = s, A_t = a]$$

Key Elements of a RL problem - The model of the environment

Stochastic versus Deterministic Environment:

- Deterministic: The outcome of each action is known (ex chess)
- Stochastic: The outcome of each action is not known (ex: dice game)

Fully Observable versus Partially Observable

- Fully Observable: The observation of the agent is the observation of the environment (ex: chess)
- Partially Observable: The agent observes indirectly the environment: the agent state is different from the environment state (ex: poker)


Key Elements of a RL problem - The model of the environment

Episodic versus continuing (non-episodic):

- Episodic: There is a terminal state (ex: game)
- Non-Episodic: There is no terminal state (ex: robotic navigation)

Discrete versus Continuous:

- Discrete: Action Space is a discrete set of actions
- (ex: Frozen Lake)
- Continuous: Action Space is continuous (ex: Mountain Car)



Part I - The Mathematical Framework for Reinforcement Learning: The Markov Decision Process

The Markov Decision Process (MDP)

A Markov Decision Process formally describes an RL environment when:

- The environment is fully observable
- The current state follows the Markov Property: $P(S_{t+1}|S_t) = P(S_{t+1}|S_t, \dots, S_1)$

A MDP is a tuple $\langle S, A, P, R, \gamma \rangle$

- Set of States S
- Set of Actions A
- Transition Probabilities P
- Reward Probabilities R
- Discount rate γ

The Markov Decision Process (MDP)

A MDP is a tuple $\langle S, A, P, R, \gamma \rangle$

- **Set of States S:** states that the agent can be in.
- **Set of Actions A:** Actions that can be performed by the agent from one state to another
- **Transition Probabilities P:** Probability of moving from one state s to another state s' by performing a :

$$p(s'|s, a) = P \left[S_{t+1} = s' \mid S_t = s, A_t = a \right]$$

- **Reward Probabilities R:** probability of reward acquired by the agent from being in one state s and performing action a :

$$r(a, s) = \mathbb{E} [r_{t+1} \mid S_t = s, A_t = a]$$

- **Discount rate γ :** controls the importance of immediate and future rewards

Value Functions in a MDP - Value Function V

How to measure how “good” is a policy?

- **Value Function $V(s)$:** measure how good is being in a state: it is the **expected return of the agent when being in state s**

$$V_{\pi}(s) = \mathbb{E}_{\pi} \left(G_t \middle| S_t = s \right) = \mathbb{E}_{\pi} \left(\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \middle| S_t = s \right)$$

Value Functions in a MDP - Action Value Function Q

How to measure how “good” is a policy?

- **Action-Value Function $Q(s,a)$:** measure how good is being in state s and performing an action a = agent’s expected return when being in state s and performing action a

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi} \left(G_t \middle| S_t = s, A_t = a \right) = \mathbb{E} \left(\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \middle| S_t = s, A_t = a, \pi \right)$$

Optimal Value Functions of a MDP

Goal of the agent = maximizing its total reward > Maximizing the value functions on all the possible policies he can follow

- Optimal Value Function v_*

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

- Optimal Action Value Function q_*

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

Optimal Policy of a MDP

Goal of the agent = maximizing its total reward > Finding the optimal policy

- Associated optimal policy π_* : a policy is said to be optimal when it dominates over any policy in every state:

$$\pi^* \in \arg \max_{\pi} V_{\pi}$$
$$\pi^* \text{ is optimal} \Leftrightarrow \forall s \in S, \forall \pi, V_{\pi^*}(s) \geq V_{\pi}(s)$$

- A MDP is solved when the optimal Value function (q_* or v_*) is reached: it is equivalent to reach the optimal policy

Bellman Equations - Bellman Expectation Equation for V

Decomposition of the Value Function between immediate reward and the Value function of the next state:

$$v_{\pi}(s) = \mathbb{E} \left[r_{t+1} + \gamma v_{\pi}(S_{t+1}) \middle| S_t = s \right]$$
$$v_{\pi}(s) = \sum_{a \in A} \pi(a|s) q_{\pi}(s, a) = \sum_{a \in A} \pi(a|s) \left(r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) v_{\pi}(s') \right)$$

Bellman Equations - Bellman Expectation Equation for Q

Decomposition of the State-Action Value Function between immediate reward and the Value function of the next state:

$$q_{\pi}(s, a) = \mathbb{E} \left[r_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) \middle| S_t = s, A_t = a \right]$$

$$q_{\pi}(s, a) = r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) \sum_{a' \in A} \pi(a'|s') q_{\pi}(a', s')$$

Bellman Optimality Equation for V

Decomposition of the Optimal Value Function between immediate reward and the Value function of the next state:

$$v_*(s) = \max_{a \in A} q_*(s, a)$$
$$v_*(s) = \max_{a \in A} \left(r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) v_*(s') \right)$$

To solve a MDP, we will solve the Bellman Optimality Equation with iterative methods.

Bellman Optimality Equation for Q

Decomposition of the Optimal State-Action Value Function between immediate reward and the State-Action Value function of the next state:

$$q_*(s, a) = r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) v_*(s')$$

$$q_*(s, a) = r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) \max_{a'} q_*(s', a')$$

To solve a MDP, we will solve the Bellman Optimality Equation with iterative methods.

Solving a MDP : Learning versus Planning

$$\text{MDP} \langle S, A, P, R, \gamma \rangle$$

- **Case 1: Full knowledge of the MDP:** the transition state matrix P and the reward probabilities R (=the dynamics of the MDP) are known.
 - > MDP solved **with Dynamic Programming Methods**: the agent performs computation with its model to improve its policy
- **Case 2: The model of the environment is unknown.**
 - > MDP solved **with Reinforcement Learning Methods**: the agent interacts with its environment to improve its policy

Part III: Solving a MDP with Dynamic Programming

Dynamic Programming

Technique to solve complex problems.

- The problem is broken into sub-problems: for each sub-problem, we compute and store the solution.
- If the same sub-problem occurs, we use the already computed solution

Dynamic Programming Algorithms to solve a MDP

- The Bellman equation allows to decompose the MDP recursively
- Value Functions allow to store and reuse function

Dynamic Programming

DP algorithms to solve a MDP

- **Value Iteration**
- **Policy Iteration**

Value Iteration Algorithm - Intuition

- Based on updates of the Value Function $V(s)$
- We start off with a random Value Function
- We improve the Value Function by an iterative application of the Bellman Optimality Equation: at each timestep k , for each state s

$$v_{k+1}(s) = \max_{a \in A} \left\{ r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) v_k(s') \right\}$$

- Theorem: the value function converges towards the optimal value function v^*

Value Iteration Algorithm

Value iteration

Initialize array V arbitrarily (e.g., $V(s) = 0$ for all $s \in \mathcal{S}^+$)

Repeat

$$\Delta \leftarrow 0$$

For each $s \in \mathcal{S}$:

$$v \leftarrow V(s)$$

$$V(s) \leftarrow \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$$

$$\Delta \leftarrow \max(\Delta, |v - V(s)|)$$

until $\Delta < \theta$ (a small positive number)

Output a deterministic policy, $\pi \approx \pi_*$, such that

$$\pi(s) = \arg \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$$

Value Iteration - Example

Value iteration

Initialize array V arbitrarily (e.g., $V(s) = 0$ for all $s \in \mathcal{S}^+$)

Repeat

$$\Delta \leftarrow 0$$

For each $s \in \mathcal{S}$:

$$v \leftarrow V(s)$$

$$V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$$

$$\Delta \leftarrow \max(\Delta, |v - V(s)|)$$


until $\Delta < \theta$ (a small positive number)

Output a deterministic policy, $\pi \approx \pi_*$, such that

$$\pi(s) = \arg \max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$$

Policy Iteration - Intuition

- Based on iterative update of the Policy π
- We start off with a random policy
- If this policy is not optimal, then the find new improved policy
- We repeat this process until we find the optimal policy
- Algorithm based on two steps:
- **Policy Evaluation:** The policy is evaluated using the value-function.
- **Policy Improvement:** Upon evaluating the value function, if it is not optimal, we find a new improved policy



Policy Iteration - Policy Evaluation and Improvement

1. Policy evaluation for current policy π_k :
 - Repeat until convergence:

$$v_{i+1}^{\pi_k}(s) \leftarrow \sum_{s' \in S} p(s'|s, a) [r(s, \pi(s), s') + \gamma v_i^{\pi_k}(s')]$$

2. Policy Improvement: find the best action according to one step look-ahead:

$$\pi_{k+1}(s) \leftarrow \arg \max_a \sum_{s' \in S} p(s'|s, a) [r(s, a, s') + \gamma v^{\pi_k}(s')]$$

Policy Iteration Algorithm

Policy iteration (using iterative policy evaluation)

1. Initialization

$V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$

2. Policy Evaluation

Repeat

$\Delta \leftarrow 0$

For each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s',r|s,\pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$ (a small positive number)

3. Policy Improvement

policy-stable $\leftarrow true$

For each $s \in \mathcal{S}$:

old-action $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

If *old-action* $\neq \pi(s)$, then *policy-stable* $\leftarrow false$

If *policy-stable*, then stop and return $V \approx v_*$ and $\pi \approx \pi_*$; else go to 2

Part IV: Reinforcement Learning with TD Learning and Monte-Carlo Methods

Model-Free Prediction and Model-Free Control

- Dynamic Programming : solve a known MDP (probability transition matrix is known)
- Reinforcement Learning solves an unknown MDP with 2 steps:
- **Model free prediction:** Estimate the value function of an unknown MDP: $V(s)$ or $Q(s,a)$
- **Model-free control:** Optimise the value function of an unknown MDP: find $V^*(s)$ or $Q^*(s,a)$

> 2 main reinforcement learning methods: Monte-Carlo and TD Learning

Monte-Carlo Reinforcement Learning

- Goal: learn v_{π} from episodes of experience under policy π by using the **empirical mean return**
- Episode= samples of (state, action, rewards) from initial state until terminal state (end of episode):
- Recall the return formula:

$$G_t = r_{t+1} + \gamma r_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

- The value function $V(s)$ is estimated by taking the average return over the episodes for each state s
- 2 variants for Monte-Carlo Policy Evaluation: First-Visit Monte-Carlo and Every-Visit Monte-Carlo

Monte-Carlo Policy Evaluation

For each state s visited during the episode:

- First-visit MC: The first-time the state is visited in the episode, increment counter state: $N(s) \leftarrow N(s) + 1$
- Every-visit MC: Every time the state s is visited in the episode, increment counter state: $N(s) \leftarrow N(s) + 1$
- Increment total return for the state: $S(s) \leftarrow S(s) + G_t$
- Value Function V estimated by mean return: $V(s) = S(s) / N(s)$
- Theorem (Law of Large Number): $V(s) \rightarrow v_{\pi(s)}$ as $N(s) \rightarrow \infty$

First-Visit MC Prediction - Algorithm

First-visit MC prediction, for estimating $V \approx v_\pi$

Initialize:

$\pi \leftarrow$ policy to be evaluated

$V \leftarrow$ an arbitrary state-value function

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

Repeat forever:

Generate an episode using π

For each state s appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s

Append G to $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

Monte-Carlo Control

- Monte-Carlo Control is based on the Policy Iteration Idea:
- **Policy Evaluation:** Learn $v \sim v_\pi$ with Monte-Carlo Policy Iteration Algorithm (instead of solving by dynamic programming)
- **Policy Improvement:** Greedy Policy Improvement
- BUT Greedy-Policy Improvement with V requires a model of the MDP:

$$\pi_{k+1}(s) \leftarrow \arg \max_a \sum_{s' \in S} p(s'|s, a) [r(s, a, s') + \gamma v^{\pi_k}(s')]$$

- Instead, Greedy-Policy Improvement with $Q(s,a)$ is model-free:

$$\pi'(s) = \arg \max_{a \in A} q(s, a)$$

> MC Control uses the State-Value function Q instead of the Value Function V

Monte-Carlo Control: From State Function $V(s)$ to State-Action Value Function Q

- Without a model, it is useful to estimate action values= the action of state-value pairs rather than state values.
- One must explicitly estimate **the value of each action** in order for the values to be useful in suggesting a policy.
- Recall the Action Value Function Q formula:

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi} \left(G_t \middle| S_t = s, A_t = a \right) = \mathbb{E} \left(\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \middle| S_t = s, A_t = a, \pi \right)$$

> Problem: how to know about the state-action value if we haven't taken this action in that state? **This is the problem of exploration**

Exploration and exploitation tradeoff

Reinforcement Learning is a trial and error process: an agent discover a good policy through experiences with the environment

- **Exploration** finds more information about the environment
- **Exploitation** exploits known information to maximise reward
- It is usually important to explore and exploit

Exploration / Exploitation trade-off examples:

- Exploitation: Go to your favorite restaurant
- Exploration: Try a new restaurant

ϵ -greedy exploration

- The simplest idea to ensure exploration of all possible actions
- ϵ is a small number: $\epsilon \in [0, 1]$
- With proba $1-\epsilon$, the greedy action is chosen
- With proba ϵ , a random action is chosen

$$\pi(a|s) = \begin{cases} 1 - \epsilon & \text{if } a^* = \arg \max_{a \in A} q(s, a) \\ \epsilon & \text{otherwise} \end{cases}$$

> **We use a epsilon-greedy improvement for policy improvement in MC Control**

Monte-Carlo Control with ϵ -greedy exploration

On-policy first-visit MC control (for ϵ -soft policies), estimates $\pi \approx \pi_*$

Initialize, for all $s \in \mathcal{S}$, $a \in \mathcal{A}(s)$:

$Q(s, a) \leftarrow$ arbitrary

$Returns(s, a) \leftarrow$ empty list

$\pi(a|s) \leftarrow$ an arbitrary ϵ -soft policy

Repeat forever:

(a) Generate an episode using π

(b) For each pair s, a appearing in the episode:

$G \leftarrow$ the return that follows the first occurrence of s, a

Append G to $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

(c) For each s in the episode:

$A^* \leftarrow \arg \max_a Q(s, a)$

(with ties broken arbitrarily)

For all $a \in \mathcal{A}(s)$:

$$\pi(a|s) \leftarrow \begin{cases} 1 - \epsilon + \epsilon/|\mathcal{A}(s)| & \text{if } a = A^* \\ \epsilon/|\mathcal{A}(s)| & \text{if } a \neq A^* \end{cases}$$

Model-Free Prediction with Temporal-Difference (TD) Learning

- TD methods learn directly from episodes of experience like MC Methods
- But **learn from incomplete episodes by bootstrapping**: update is done every timestep and not after each episode
- TD updates a guess towards a guess.
- The Value $V(S_t)$ is updated toward estimated return: the immediate reward + the estimated reward for the rest of the trajectory:

$$R_{t+1} + \gamma V(S_{t+1})$$

Model-Free Prediction with Temporal-Difference (TD) Learning - The TD update

- The Value $V(S_t)$ is updated toward estimated return: the immediate reward + the estimated reward for the rest of the trajectory $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$


$R_{t+1} + \gamma V(S_{t+1})$ is the TD-update

$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is the TD-error

TD Control

- Same idea than to MC Control: Apply TD Prediction to $Q(s,a)$
- Use a epsilon-greedy improvement
- Update every time-step instead of every episode for Monte-Carlo
- **2 TD learning variants**
- SARSA: On-Policy Control
- Q-Learning: Off-Policy Control

SARSA update



A' is chosen following
a policy

**=> On-policy
algorithm**

$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A))$$

SARSA algorithm

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Repeat (for each step of episode):

Take action A , observe R, S'


Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal

Q-learning update



$$Q(S, A) \leftarrow Q(S, A) + \alpha \left(R + \gamma \max_{a'} Q(S', a') - Q(S, A) \right)$$

Q-learning algorithm

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Repeat (for each step of episode):

Choose A from S using policy derived from Q (e.g., ϵ -greedy)

Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$S \leftarrow S'$;

until S is terminal

Monte-Carlo versus TD-Learning

TD can learn before knowing the final outcome:

- TD learns online at every timestep
- MC must wait until the end of the episode before return is known

TD can learn without knowing the final outcome

- TD can learn from incomplete sequences
- TD can learn in continuing environments
- MC can only learn in episodic environments

MC versus TD : Bias/Variance trade-off

- In MC: Return G_t is **unbiased estimate** of $V_{\pi}(s)$:

$$G_t = r_{t+1} + \gamma r_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

- In TD: TD target is **biased estimate** of $V_{\pi}(s)$:

$$R_{t+1} + \gamma V(S_{t+1})$$

TD target is much lower variance than the Return:

- Return depends on many random samples of actions, transitions, rewards
- TD update depends on one random action, transition, reward

MC versus TD: advantages and disadvantages

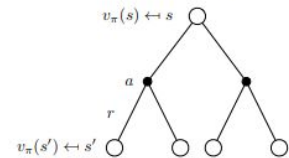

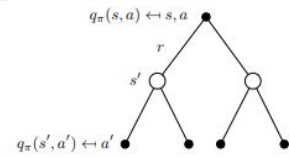
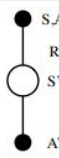
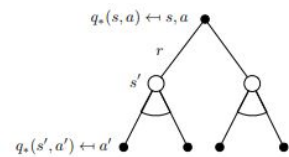
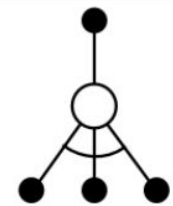
MC has high variance and zero bias:

- Not very sensitive to initial value of state
- Very simple to understand and use
- Only useable in episodic environments

TD has low variance, some bias:

- Usually more sample-efficient than MC (converge with few numbers of episodes)
- More sensitive to initial value of state

Summary: DP versus TD-learning

	Full Backup (DP)	Sample Backup (TD)
Bellman Expectation Equation for $v_{\pi}(s)$	 <p>Iterative Policy Evaluation</p>	 <p>TD Learning</p>
Bellman Expectation Equation for $q_{\pi}(s, a)$	 <p>Q-Policy Iteration</p>	 <p>Sarsa</p>
Bellman Optimality Equation for $q_{*}(s, a)$	 <p>Q-Value Iteration</p>	 <p>Q-Learning</p>

Summary: DP versus TD-learning

<i>Full Backup (DP)</i>	<i>Sample Backup (TD)</i>
Iterative Policy Evaluation $V(s) \leftarrow \mathbb{E}[R + \gamma V(S') \mid s]$	TD Learning $V(S) \stackrel{\alpha}{\leftarrow} R + \gamma V(S')$
Q-Policy Iteration $Q(s, a) \leftarrow \mathbb{E}[R + \gamma Q(S', A') \mid s, a]$	Sarsa $Q(S, A) \stackrel{\alpha}{\leftarrow} R + \gamma Q(S', A')$
Q-Value Iteration $Q(s, a) \leftarrow \mathbb{E}\left[R + \gamma \max_{a' \in \mathcal{A}} Q(S', a') \mid s, a\right]$	Q-Learning $Q(S, A) \stackrel{\alpha}{\leftarrow} R + \gamma \max_{a' \in \mathcal{A}} Q(S', a')$

Monte-Carlo Exploration Starts Algorithm for Q Evaluation

- Same concept than MC for state values except than we talk about **visits of state-actions pairs (s,a)**
- **To ensure exploration of all (s,a) pairs, we introduce the exploring starts:** we make each episode to start in a specific (s,a) pairs with every pair having a non-zero probability of being selected as the start.